


# spaCy Lib for NLP

- Gain Meaning from Unstructured Text



# **Chapter 1: Finding words, phrases, names and concepts**

# The nlp object

```
# Import the English language class  
from spacy.lang.en import English
```

```
# Create the nlp object  
nlp = English()
```

- contains the processing pipeline
- includes language-specific rules for tokenization etc.


# The Doc object

```
# Created by processing a string of text with the nlp object  
doc = nlp("Hello world!")
```

```
# Iterate over tokens in a Doc  
for token in doc:  
    print(token.text)
```

```
Hello  
world  
!
```

# The Token object



```
doc = nlp("Hello world!")
```

```
# Index into the Doc to get a single Token
```

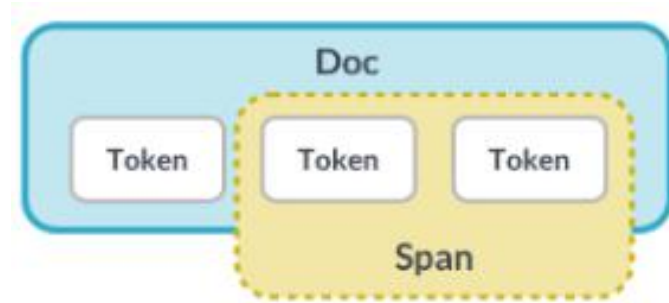
```
token = doc[1]
```

```
# Get the token text via the .text attribute
```

```
print(token.text)
```

```
world
```

# The Span object



```
doc = nlp("Hello world!")
```

```
# Index into the Doc to get a single Token  
token = doc[1]
```

```
# Get the token text via the .text attribute  
print(token.text)  
world
```

# Lexical Attributes

```
doc = nlp("It costs $5.")
```

```
print("Index:  ", [token.i for token in doc])
```

```
print("Text:    ", [token.text for token in doc])
```

```
print("is_alpha:", [token.is_alpha for token in doc])
```

```
print("is_punct:", [token.is_punct for token in doc])
```

```
print("like_num:", [token.like_num for token in doc])
```

```
Index:  [0, 1, 2, 3, 4]
```

```
Text:    ['It', 'costs', '$', '5', '.']
```

```
is_alpha: [True, True, False, False, False]
```

```
is_punct: [False, False, False, False, True]
```

```
like_num: [False, False, False, True, False]
```

# Statistical models



# What are statistical models?

- Enable spaCy to predict linguistic attributes in context
  - Part-of-speech tags
  - Syntactic dependencies
  - Named entities
- Trained on labeled example texts
- Can be updated with more examples to fine-tune predictions

# Model Packages

```
$ python -m spacy download en_core_web_sm  
import spacy
```

```
nlp = spacy.load("en_core_web_sm")
```

- Binary weights
- Vocabulary
- Meta information (language, pipeline)



# Predicting Part-of-speech Tags

```
import spacy

# Load the small English model
nlp = spacy.load("en_core_web_sm")

# Process a text
doc = nlp("She ate the pizza")

# Iterate over the tokens
for token in doc:
    # Print the text and the predicted part-of-speech tag
    print(token.text, token.pos_)
```

```
She PRON
ate VERB
the DET
pizza NOUN
```


# Predicting Syntactic Dependencies

for token in doc:

```
print(token.text, token.pos_, token.dep_, token.head.text)
```

She PRON nsubj ate  
ate VERB ROOT ate  
the DET det pizza  
pizza NOUN dobj ate

# Dependency label scheme



Label	Description	Example
nsubj	nominal subject	She
dobj	direct object	pizza
det	determiner (article)	the

# Predicting Named Entities

Apple ORG is looking at buying U.K. GPE startup for \$1 billion MONEY

# Process a text

```
doc = nlp("Apple is looking at buying U.K. startup for $1 billion")
```

# Iterate over the predicted entities

for ent in doc.ents:

    # Print the entity text and its label

```
    print(ent.text, ent.label_)
```

Apple ORG

U.K. GPE

\$1 billion MONEY

## Tip: the `spacy.explain` method

Get quick definitions of the most common tags and labels.

```
spacy.explain("GPE")  
'Countries, cities, states'
```

```
spacy.explain("NNP")  
'noun, proper singular'
```

```
spacy.explain("dobj")  
'direct object'
```

# Rule-based matching



# Why not just regular expressions?

- Match on Doc objects, not just strings
- Match on tokens and token attributes
- Use the model's predictions
- Example: "duck" (verb) vs. "duck" (noun)

# Match patterns

- Lists of dictionaries, one per token
- Match exact token texts  
[{"TEXT": "iPhone"}, {"TEXT": "X"}]
- Match lexical attributes  
[{"LOWER": "iphone"}, {"LOWER": "x"}]
- Match any token attributes  
[{"LEMMA": "buy"}, {"POS": "NOUN"}]

# Using the Matcher (1)

```
import spacy
# Import the Matcher
from spacy.matcher import Matcher
# Load a model and create the nlp object
nlp = spacy.load("en_core_web_sm")
# Initialize the matcher with the shared vocab
matcher = Matcher(nlp.vocab)
# Add the pattern to the matcher
pattern = [{"TEXT": "iPhone"}, {"TEXT": "X"}]
matcher.add("IPHONE_PATTERN", None, pattern)
# Process some text
doc = nlp("Upcoming iPhone X release date leaked")
# Call the matcher on the doc
matches = matcher(doc)
```

## Using the Matcher (2)

```
# Call the matcher on the doc
```

```
doc = nlp("Upcoming iPhone X release date leaked")
```

```
matches = matcher(doc)
```

```
# Iterate over the matches
```

```
for match_id, start, end in matches:
```

```
    # Get the matched span
```

```
    matched_span = doc[start:end]
```

```
    print(matched_span.text)
```

iPhone X

match\_id: hash value of the pattern name

start: start index of matched span

end: end index of matched span

# Matching lexical attributes

```
pattern = [  
    {"IS_DIGIT": True},  
    {"LOWER": "fifa"},  
    {"LOWER": "world"},  
    {"LOWER": "cup"},  
    {"IS_PUNCT": True}  
]  
doc = nlp("2018 FIFA World Cup: France won!")
```

2018 FIFA World Cup:

# Matching other token attributes

```
pattern = [  
    {"LEMMA": "love", "POS": "VERB"},  
    {"POS": "NOUN"}  
]  
doc = nlp("I loved dogs but now I love cats more.")
```

loved dogs

love cats

# Using operators and quantifiers (1)

```
pattern = [  
    {"LEMMA": "buy"},  
    {"POS": "DET", "OP": "?"}, # optional: match 0 or 1 times  
    {"POS": "NOUN"}  
]  
doc = nlp("I bought a smartphone. Now I'm buying apps.")
```

bought a smartphone

buying apps

# Using operators and quantifiers (2)

Example	Description
<code>{"OP": "!"}</code>	Negation: match 0 times
<code>{"OP": "?"}</code>	Optional: match 0 or 1 times
<code>{"OP": "+"}</code>	Match 1 or more times
<code>{"OP": "*"}</code>	Match 0 or more times



```
import spacy
from spacy.matcher import Matcher
nlp = spacy.load("en_core_web_sm")
matcher = Matcher(nlp.vocab)

doc = nlp(
    "After making the iOS update you won't notice a radical system-wide "
    "redesign: nothing like the aesthetic upheaval we got with iOS 7. Most of "
    "iOS 11's furniture remains the same as in iOS 10. But you will discover "
    "some tweaks once you delve a little deeper.")

# Write a pattern for full iOS versions ("iOS 7", "iOS 11", "iOS 10")
pattern = [{"TEXT": "iOS"}, {"IS_DIGIT": True}]
# Add the pattern to the matcher and apply the matcher to the doc
matcher.add("IOS_VERSION_PATTERN", None, pattern)
matches = matcher(doc)
print("Total matches found:", len(matches))
# Iterate over the matches and print the span text
for match_id, start, end in matches:
    print("Match found:", doc[start:end].text)
```

```
import spacy
from spacy.matcher import Matcher
nlp = spacy.load("en_core_web_sm")
matcher = Matcher(nlp.vocab)

doc = nlp(
    "i downloaded Fortnite on my laptop and can't open the game at all. Help? "
    "so when I was downloading Minecraft, I got the Windows version where it "
    "is the '.zip' folder and I used the default program to unpack it... do "
    "I also need to download Winzip?")

# Write a pattern that matches a form of "download" plus proper noun
pattern = [{"LEMMA": "download"}, {"POS": "PROPN"}]
# Add the pattern to the matcher and apply the matcher to the doc
matcher.add("DOWNLOAD_THINGS_PATTERN", None, pattern)
matches = matcher(doc)
print("Total matches found:", len(matches))
# Iterate over the matches and print the span text
for match_id, start, end in matches:
    print("Match found:", doc[start:end].text)
```

```
import spacy
from spacy.matcher import Matcher
nlp = spacy.load("en_core_web_sm")
matcher = Matcher(nlp.vocab)

doc = nlp(
    "Features of the app include a beautiful design, smart search, automatic "
    "labels and optional voice responses.")

# Write a pattern for adjective plus one or two nouns
pattern = [{"POS": _____}, {"POS": _____}, {"POS": _____, "OP": _____}]
# Add the pattern to the matcher and apply the matcher to the doc
matcher.add("ADJ_NOUN_PATTERN", None, pattern)
matches = matcher(doc)
print("Total matches found:", len(matches))

# Iterate over the matches and print the span text
for match_id, start, end in matches:
    print("Match found:", doc[start:end].text)
```

## **Chapter 2: Large-scale data analysis with spaCy**

# **Data Structures (1): Vocab, Lexemes and StringStore**

# Shared vocab and string store (1)

- Vocab: stores data shared across multiple documents
- To save memory, spaCy encodes all strings to hash values
- Strings are only stored once in the StringStore via `nlp.vocab.strings`
- String store: lookup table in both directions

```
coffee_hash = nlp.vocab.strings["coffee"]
```

```
coffee_string = nlp.vocab.strings[coffee_hash]
```

- Hashes can't be reversed – that's why we need to provide the shared vocab

```
# Raises an error if we haven't seen the string before
```

```
string = nlp.vocab.strings[3197928453018144401]
```

## Shared vocab and string store (2)

- Look up the string and hash in `nlp.vocab.strings`

```
doc = nlp("I love coffee")
```

```
print("hash value:", nlp.vocab.strings["coffee"])
```

```
print("string value:", nlp.vocab.strings[3197928453018144401])
```

```
hash value: 3197928453018144401
```

```
string value: coffee
```

- The doc also exposes the vocab and strings

```
doc = nlp("I love coffee")
```

```
print("hash value:", doc.vocab.strings["coffee"])
```

```
hash value: 3197928453018144401
```

# Lexemes: entries in the vocabulary

- A Lexeme object is an entry in the vocabulary

```
doc = nlp("I love coffee")
```

```
lexeme = nlp.vocab["coffee"]
```

```
# Print the lexical attributes
```

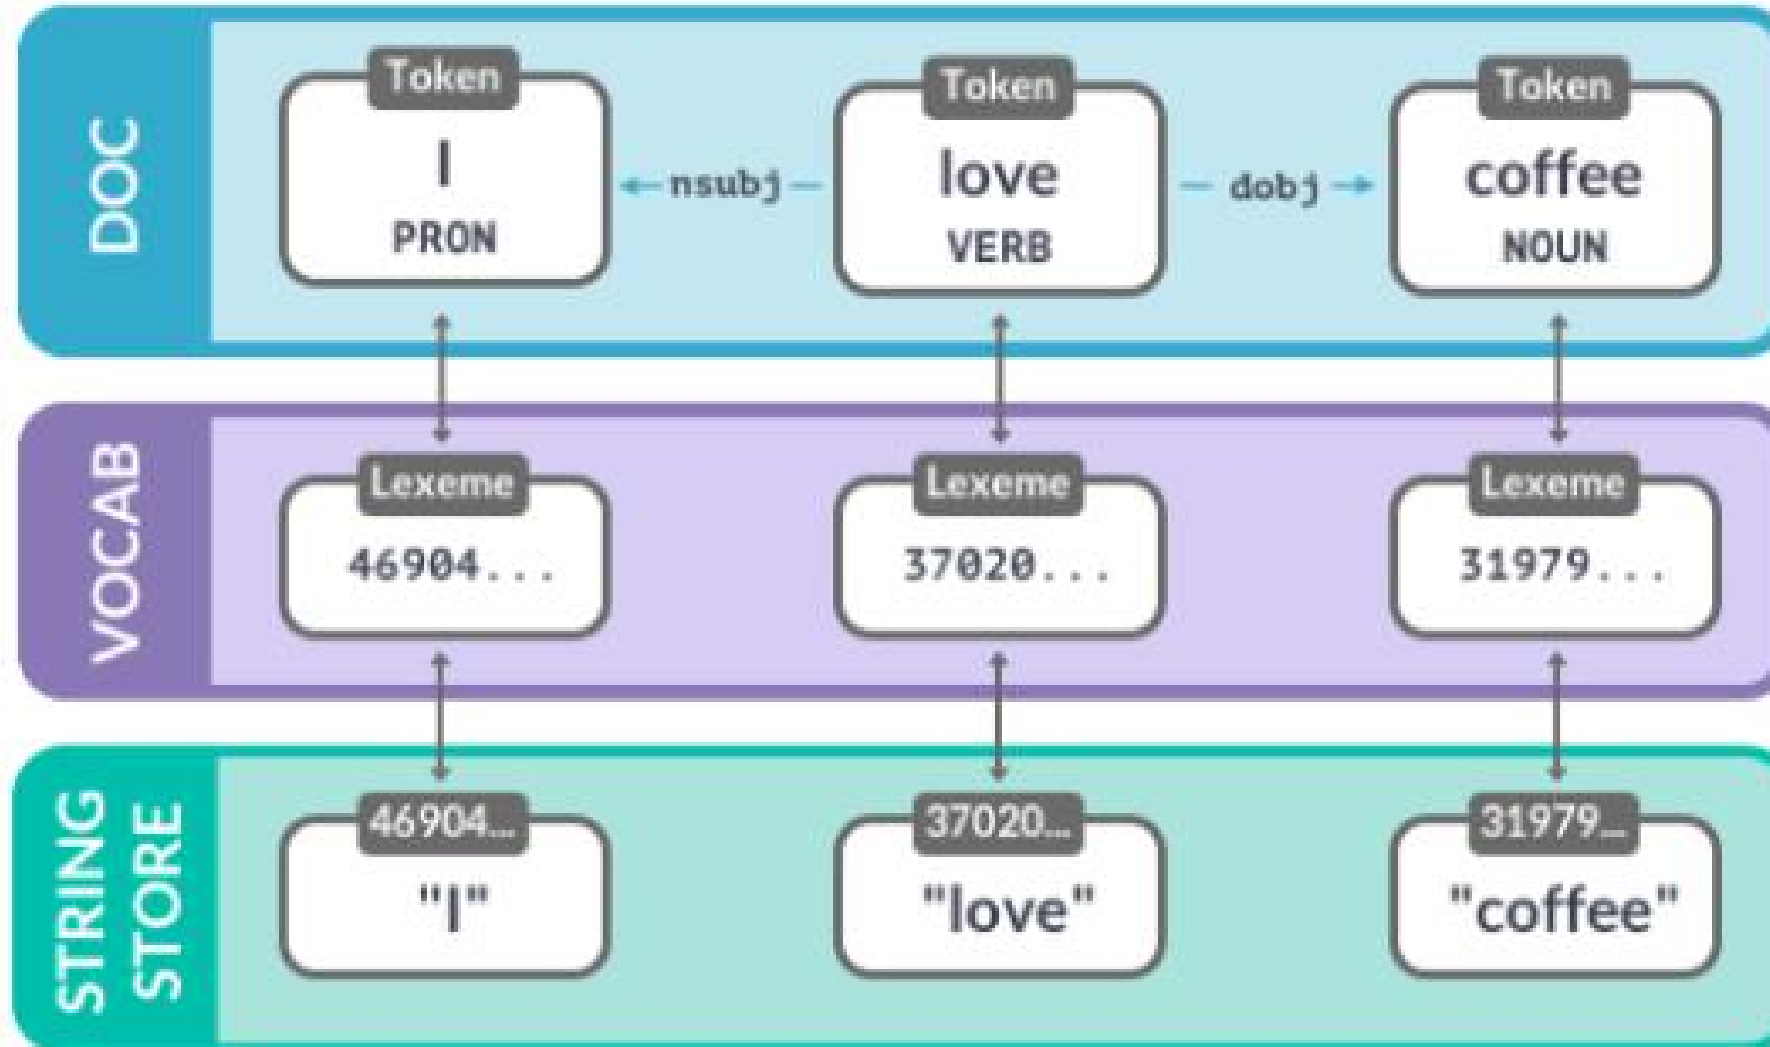
```
print(lexeme.text, lexeme.orth, lexeme.is_alpha)
```

```
coffee 3197928453018144401 True
```

- Contains the context-independent information about a word
  - Word text: `lexeme.text` and `lexeme.orth` (the hash)
  - Lexical attributes like `lexeme.is_alpha`
  - Not context-dependent part-of-speech tags, dependencies or entity labels



# Vocab, hashes and lexemes



# Vocab, hashes and lexemes

```
from spacy.lang.en import English
```

```
nlp = English()
```

```
doc = nlp("I have a cat")
```

```
# Look up the hash for the word "cat"
```

```
cat_hash = _____._____._____
```

```
print(cat_hash)
```

```
# Look up the cat_hash to get the string
```

```
cat_string = _____._____._____
```

```
print(cat_string)
```

# Vocab, hashes and lexemes

```
import spacy
```

```
nlp = spacy.load("en_core_web_sm")  
doc = nlp("David Bowie is a PERSON")
```

```
# Look up the hash for the string label "PERSON"  
person_hash = nlp.vocab.strings["PERSON"]  
print(person_hash)
```

```
# Look up the person_hash to get the string  
person_string = nlp.vocab.strings[person_hash]  
print(person_string)
```

## **Data Structures (2): Doc, Span and Token**

# The Doc object

# Create an nlp object

```
from spacy.lang.en import English
```

```
nlp = English()
```

# Import the Doc class

```
from spacy.tokens import Doc
```

# The words and spaces to create the doc from

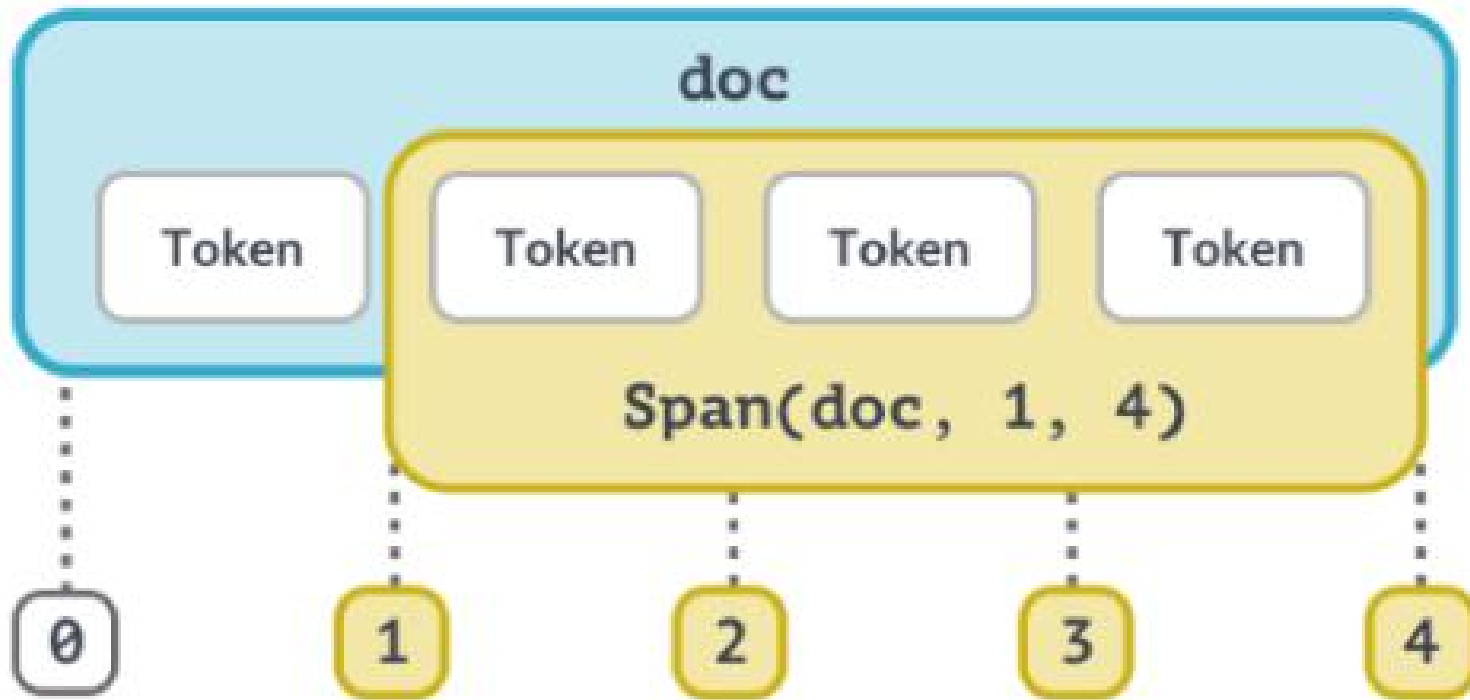
```
words = ["Hello", "world", "!"]
```

```
spaces = [True, False, False]
```

# Create a doc manually

```
doc = Doc(nlp.vocab, words=words, spaces=spaces)
```

# The Span object (1)



# The Span object (2)

```
# Import the Doc and Span classes  
from spacy.tokens import Doc, Span
```

```
# The words and spaces to create the doc from  
words = ["Hello", "world", "!"]  
spaces = [True, False, False]
```

```
# Create a doc manually  
doc = Doc(nlp.vocab, words=words, spaces=spaces)
```

```
# Create a span manually  
span = Span(doc, 0, 2)
```

```
# Create a span with a label  
span_with_label = Span(doc, 0, 2, label="GREETING")
```

```
# Add span to the doc.ents  
doc.ents = [span_with_label]
```

## Best practices

- Doc and Span are very powerful and hold references and relationships of words and sentences
  - Convert result to strings as late as possible
  - Use token attributes if available – for example, `token.i` for the token index
- Don't forget to pass in the shared vocab



# Best practices

```
import spacy
```

```
nlp = spacy.load("en_core_web_sm")
```

```
doc = nlp("Berlin is a nice city")
```

```
# Iterate over the tokens
```

```
for token in doc:
```

```
    # Check if the current token is a proper noun
```

```
    if token.pos_ == "PROPN":
```

```
        # Check if the next token is a verb
```

```
        if doc[token.i + 1].pos_ == "VERB":
```

```
            print("Found proper noun before a verb:", token.text)
```

# **Word vectors and semantic similarity**

# Comparing semantic similarity

- spaCy can compare two objects and predict similarity
- `Doc.similarity()`, `Span.similarity()` and `Token.similarity()`
- Take another object and return a similarity score (0 to 1)
- Important: needs a model that has word vectors included, for example:
  - ✓ `en_core_web_md` (medium model)
  - ✓ `en_core_web_lg` (large model)
  - ✗ NOT `en_core_web_sm` (small model)

# Similarity examples (1)

# Load a larger model with vectors

```
nlp = spacy.load("en_core_web_md")
```

# Compare two documents

```
doc1 = nlp("I like fast food")
```

```
doc2 = nlp("I like pizza")
```

```
print(doc1.similarity(doc2))
```

*0.8627204117787385*

# Compare two tokens

```
doc = nlp("I like pizza and pasta")
```

```
token1 = doc[2]
```

```
token2 = doc[4]
```

```
print(token1.similarity(token2))
```

*0.7369546*

## Similarity examples (2)

# Compare a document with a token

```
doc = nlp("I like pizza")
```

```
token = nlp("soap")[0]
```

```
print(doc.similarity(token))
```

```
0.32531983166759537
```

# Compare a span with a document

```
span = nlp("I like pizza and pasta")[2:5]
```

```
doc = nlp("McDonalds sells burgers")
```

```
print(span.similarity(doc))
```

```
0.619909235817623
```

# How does spaCy predict similarity?

- Similarity is determined using word vectors
- Multi-dimensional meaning representations of words
- Generated using an algorithm like Word2Vec and lots of text
- Can be added to spaCy's statistical models
- Default: cosine similarity, but can be adjusted
- Doc and Span vectors default to average of token vectors
- Short phrases are better than long documents with many irrelevant words

## Word vectors in spaCy

# Load a larger model with vectors

```
nlp = spacy.load("en_core_web_md")
```

```
doc = nlp("I have a banana")
```

# Access the vector via the token.vector attribute

```
print(doc[3].vector)
```

```
[2.02280000e-01, -7.66180009e-02,  3.70319992e-01,  
 3.28450017e-02, -4.19569999e-01,  7.20689967e-02,  
-3.74760002e-01,  5.74599989e-02, -1.24009997e-02,  
 5.29489994e-01, -5.23800015e-01, -1.97710007e-01,  
 ...
```

## Similarity depends on the application context

- Useful for many applications: recommendation systems, flagging duplicates etc.
- There's no objective definition of "similarity"
- Depends on the context and what application needs to do

```
doc1 = nlp("I like cats")
```

```
doc2 = nlp("I hate cats")
```

```
print(doc1.similarity(doc2))
```

```
0.9501447503553421
```



# Inspecting word vectors

```
import spacy
```

```
# Load the en_core_web_md model  
nlp = spacy.load("en_core_web_md")
```

```
# Process a text  
doc = nlp("Two bananas in pyjamas")
```

```
# Get the vector for the token "bananas"  
bananas_vector = doc[1].vector  
print(bananas_vector)
```

# Combining models and rules

# Statistical predictions vs. rules

	Statistical models	Rule-based systems
Use cases	application needs to <i>generalize</i> based on examples	dictionary with finite number of examples
Real-world examples	product names, person names, subject/object relationships	countries of the world, cities, drug names, dog breeds
spaCy features	entity recognizer, dependency parser, part-of-speech tagger	tokenizer, <code>Matcher</code> , <code>PhraseMatcher</code>

## Recap: Rule-based Matching

```
# Initialize with the shared vocab
```

```
from spacy.matcher import Matcher
```

```
matcher = Matcher(nlp.vocab)
```

```
# Patterns are lists of dictionaries describing the tokens
```

```
pattern = [{"LEMMA": "love", "POS": "VERB"}, {"LOWER": "cats"}]
```

```
matcher.add("LOVE_CATS", None, pattern)
```

```
# Operators can specify how often a token should be matched
```

```
pattern = [{"TEXT": "very", "OP": "+"}, {"TEXT": "happy"}]
```

```
matcher.add("VERY_HAPPY", None, pattern)
```

```
# Calling matcher on doc returns list of (match_id, start, end) tuples
```

```
doc = nlp("I love cats and I'm very very happy")
```

```
matches = matcher(doc)
```

# Adding statistical predictions

```
matcher = Matcher(nlp.vocab)
matcher.add("DOG", None, [{"LOWER": "golden"}, {"LOWER": "retriever"}])
doc = nlp("I have a Golden Retriever")
```

```
for match_id, start, end in matcher(doc):
    span = doc[start:end]
    print("Matched span:", span.text)
    # Get the span's root token and root head token
    print("Root token:", span.root.text)
    print("Root head token:", span.root.head.text)
    # Get the previous token and its POS tag
    print("Previous token:", doc[start - 1].text, doc[start - 1].pos_)
```

Matched span: Golden Retriever

Root token: Retriever

Root head token: have

Previous token: a DET

## Efficient phrase matching (1)

- PhraseMatcher like regular expressions or keyword search – but with access to the tokens!
- Takes Doc object as patterns
- More efficient and faster than the Matcher
- Great for matching large word lists

## Efficient phrase matching (2)

```
from spacy.matcher import PhraseMatcher
```

```
matcher = PhraseMatcher(nlp.vocab)
```

```
pattern = nlp("Golden Retriever")
```

```
matcher.add("DOG", None, pattern)
```

```
doc = nlp("I have a Golden Retriever")
```

```
# Iterate over the matches
```

```
for match_id, start, end in matcher(doc):
```

```
    # Get the matched span
```

```
    span = doc[start:end]
```

```
    print("Matched span:", span.text)
```

```
Matched span: Golden Retriever
```

```
import json
from spacy.lang.en import English
with open("exercises/en/countries.json") as f:
    COUNTRIES = json.loads(f.read())

nlp = English()
doc = nlp("Czech Republic may help Slovakia protect its airspace")

# Import the PhraseMatcher and initialize it
from spacy.matcher import PhraseMatcher
matcher = PhraseMatcher(nlp.vocab)

# Create pattern Doc objects and add them to the matcher
# This is the faster version of: [nlp(country) for country in COUNTRIES]
patterns = list(nlp.pipe(COUNTRIES))
matcher.add("COUNTRY", None, *patterns)

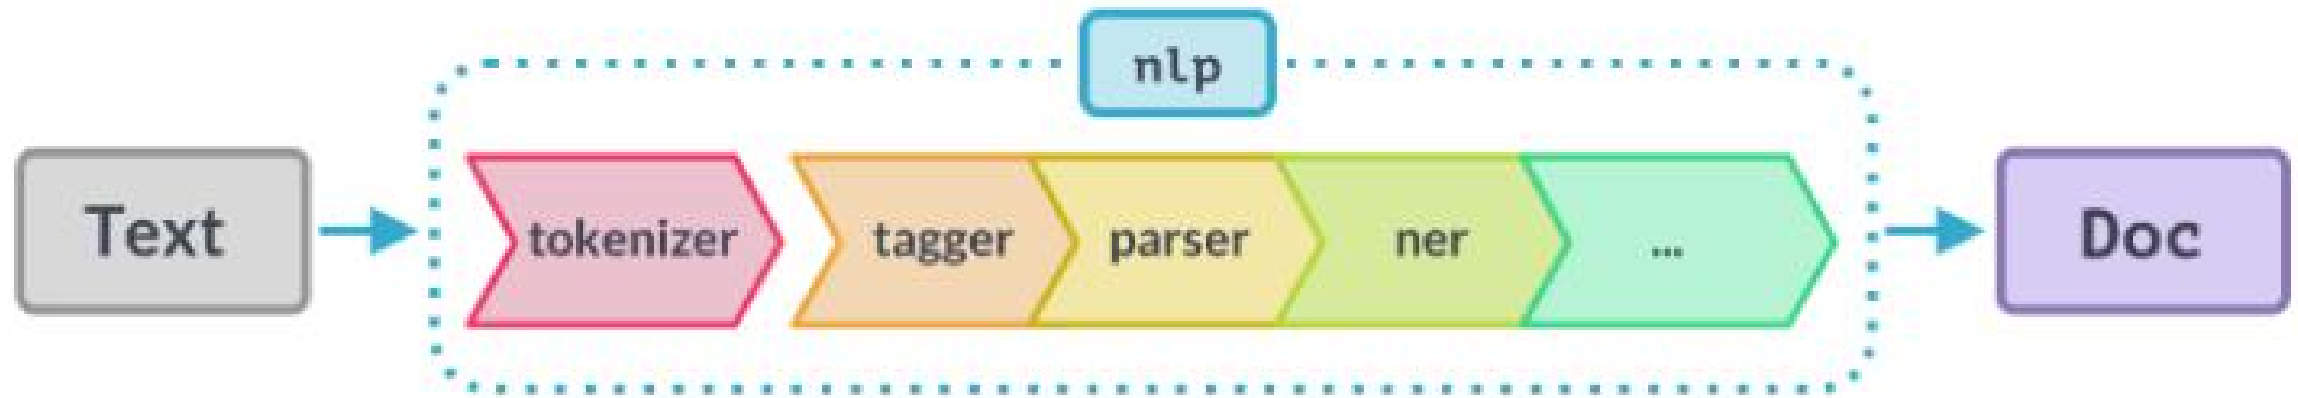
# Call the matcher on the test document and print the result
matches = matcher(doc)
print([doc[start:end] for match_id, start, end in matches])
```



## **Chapter 3: Processing Pipelines**

# Processing pipelines

# What happens when you call nlp?



```
doc = nlp("This is a sentence.")
```

# Built-in pipeline components

Name	Description	Creates
tagger	Part-of-speech tagger	<code>Token.tag</code> , <code>Token.pos</code>
parser	Dependency parser	<code>Token.dep</code> , <code>Token.head</code> , <code>Doc.sents</code> , <code>Doc.noun_chunks</code>
ner	Named entity recognizer	<code>Doc.ents</code> , <code>Token.ent_iob</code> , <code>Token.ent_type</code>
textcat	Text classifier	<code>Doc.cats</code>

# Under the hood



meta.json

```
{  
  "lang": "en",  
  "name": "core_web_sm",  
  "pipeline": ["tagger", "parser", "ner"]  
}
```

- Pipeline defined in model's meta.json in order
- Built-in components need binary data to make predictions

## Pipeline attributes

- `nlp.pipe_names`: list of pipeline component names

```
print(nlp.pipe_names)
```

```
['tagger', 'parser', 'ner']
```

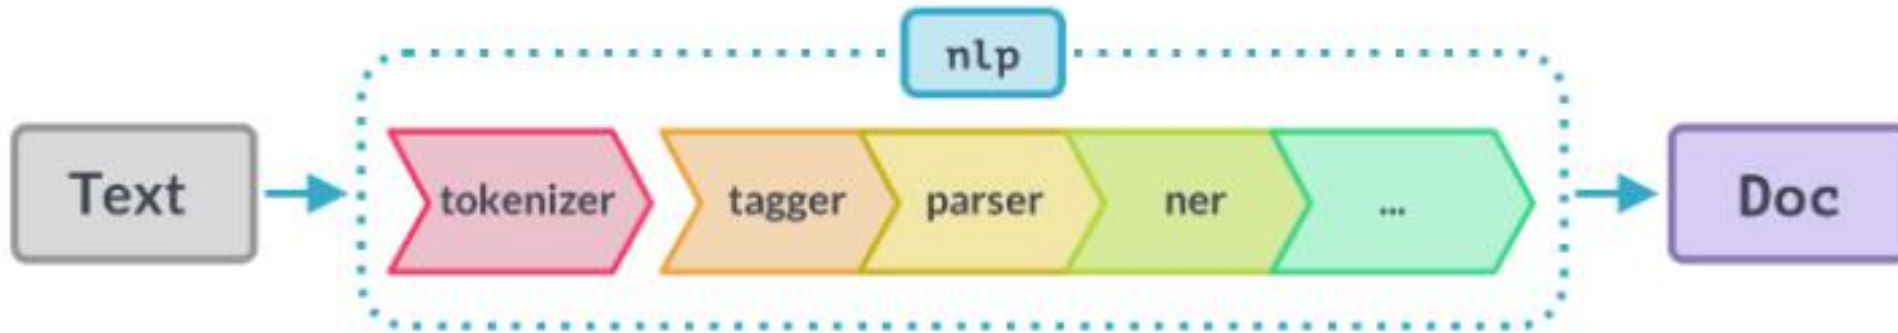
- `nlp.pipeline`: list of (name, component) tuples

```
print(nlp.pipeline)
```

```
[('tagger', <spacy.pipeline.Tagger>),  
 ('parser', <spacy.pipeline.DependencyParser>),  
 ('ner', <spacy.pipeline.EntityRecognizer>)]
```

# Custom pipeline components

## Why custom components?



- Make a function execute automatically when you call `nlp`
- Add your own metadata to documents and tokens
- Updating built-in attributes like `doc.ents`



## Anatomy of a component (1)

- Function that takes a doc, modifies it and returns it
- Can be added using the `nlp.add_pipe` method

```
def custom_component(doc):  
    # Do something to the doc here  
    return doc
```

```
nlp.add_pipe(custom_component)
```

## Anatomy of a component (2)

```
def custom_component(doc):  
    # Do something to the doc here  
    return doc
```

```
nlp.add_pipe(custom_component)
```

Argument	Description	Example
last	If True, add last	<code>nlp.add_pipe(component, last=True)</code>
first	If True, add first	<code>nlp.add_pipe(component, first=True)</code>
before	Add before component	<code>nlp.add_pipe(component, before="ner")</code>
after	Add after component	<code>nlp.add_pipe(component, after="tagger")</code>

## Example: a simple component (1)

```
# Create the nlp object
```

```
nlp = spacy.load("en_core_web_sm")
```

```
# Define a custom component
```

```
def custom_component(doc):
```

```
    # Print the doc's length
```

```
    print("Doc length:", len(doc))
```

```
    # Return the doc object
```

```
    return doc
```

```
# Add the component first in the pipeline
```

```
nlp.add_pipe(custom_component, first=True)
```

```
# Print the pipeline component names
```

```
print("Pipeline:", nlp.pipe_names)
```

```
Pipeline: ['custom_component', 'tagger', 'parser', 'ner']
```

## Example: a simple component (2)

```
# Create the nlp object
nlp = spacy.load("en_core_web_sm")

# Define a custom component
def custom_component(doc):
    # Print the doc's length
    print("Doc length:", len(doc))

    # Return the doc object
    return doc

# Add the component first in the pipeline
nlp.add_pipe(custom_component, first=True)

# Process a text
doc = nlp("Hello world!")
Doc length: 3
```

## Extension attributes

# Setting custom attributes

- Add custom metadata to documents, tokens and spans
- Accessible via the `._` property

```
doc._.title = "My document"
```

```
token._.is_color = True
```

```
span._.has_color = False
```

- Registered on the global Doc, Token or Span using the `set_extension` method

```
# Import global classes
```

```
from spacy.tokens import Doc, Token, Span
```

```
# Set extensions on the Doc, Token and Span
```

```
Doc.set_extension("title", default=None)
```

```
Token.set_extension("is_color", default=False)
```

```
Span.set_extension("has_color", default=False)
```

# Extension attribute types

1. Attribute extensions
2. Property extensions
3. Method extensions

## Attribute extensions

- Set a default value that can be overwritten

```
from spacy.tokens import Token
```

```
# Set extension on the Token with default value  
Token.set_extension("is_color", default=False)
```

```
doc = nlp("The sky is blue.")
```

```
# Overwrite extension attribute value  
doc[3]._.is_color = True
```



# Property extensions (1)

- Define a getter and an optional setter function
- Getter only called when you retrieve the attribute value

from spacy.tokens import Token

# Define getter function

```
def get_is_color(token):  
    colors = ["red", "yellow", "blue"]  
    return token.text in colors
```

# Set extension on the Token with getter

```
Token.set_extension("is_color", getter=get_is_color)
```

```
doc = nlp("The sky is blue.")  
print(doc[3]._.is_color, "-", doc[3].text)
```

*True - blue*

## Property extensions (2)

- Span extensions should almost always use a getter  
from `spacy.tokens` import `Span`

```
# Define getter function
```

```
def get_has_color(span):  
    colors = ["red", "yellow", "blue"]  
    return any(token.text in colors for token in span)
```

```
# Set extension on the Span with getter
```

```
Span.set_extension("has_color", getter=get_has_color)
```

```
doc = nlp("The sky is blue.")  
print(doc[1:4]._.has_color, "-", doc[1:4].text)  
print(doc[0:2]._.has_color, "-", doc[0:2].text)
```

*True - sky is blue*

*False - The sky*

# Method extensions

- Assign a function that becomes available as an object method
  - Lets you pass arguments to the extension function
- from spacy.tokens import Doc

# Define method with arguments

```
def has_token(doc, token_text):  
    in_doc = token_text in [token.text for token in doc]  
    return in_doc
```

# Set extension on the Doc with method

```
Doc.set_extension("has_token", method=has_token)
```

```
doc = nlp("The sky is blue.")  
print(doc._.has_token("blue"), "- blue")  
print(doc._.has_token("cloud"), "- cloud")
```

*True - blue*

*False - cloud*

```
import spacy
from spacy.tokens import Span
nlp = spacy.load("en_core_web_sm")

def get_wikipedia_url(span):
    # Get a Wikipedia URL if the span has one of the labels
    if span.label_ in ("PERSON", "ORG", "GPE", "LOCATION"):
        entity_text = span.text.replace(" ", "_")
        return "https://en.wikipedia.org/w/index.php?search=" + entity_text

# Set the Span extension wikipedia_url using get getter get_wikipedia_url
Span.set_extension("wikipedia_url", getter=get_wikipedia_url)

doc = nlp(
    "In over fifty years from his very first recordings right through to his "
    "last album, David Bowie was at the vanguard of contemporary culture."
)
for ent in doc.ents:
    # Print the text and Wikipedia URL of the entity
    print(ent.text, ent._.wikipedia_url)
```

# Scaling and performance

## Processing large volumes of text

- Use `nlp.pipe` method
- Processes texts as a stream, yields Doc objects
- Much faster than calling `nlp` on each text

BAD:

```
docs = [nlp(text) for text in LOTS_OF_TEXTS]
```

GOOD:

```
docs = list(nlp.pipe(LOTS_OF_TEXTS))
```

## Passing in context (1)

- Setting `as_tuples=True` on `nlp.pipe` lets you pass in (text, context) tuples
- Yields (doc, context) tuples
- Useful for associating metadata with the doc

```
data = [  
    ("This is a text", {"id": 1, "page_number": 15}),  
    ("And another text", {"id": 2, "page_number": 16}),  
]
```

```
for doc, context in nlp.pipe(data, as_tuples=True):  
    print(doc.text, context["page_number"])
```

*This is a text 15*

*And another text 16*

## Passing in context (2)

```
from spacy.tokens import Doc
```

```
Doc.set_extension("id", default=None)
```

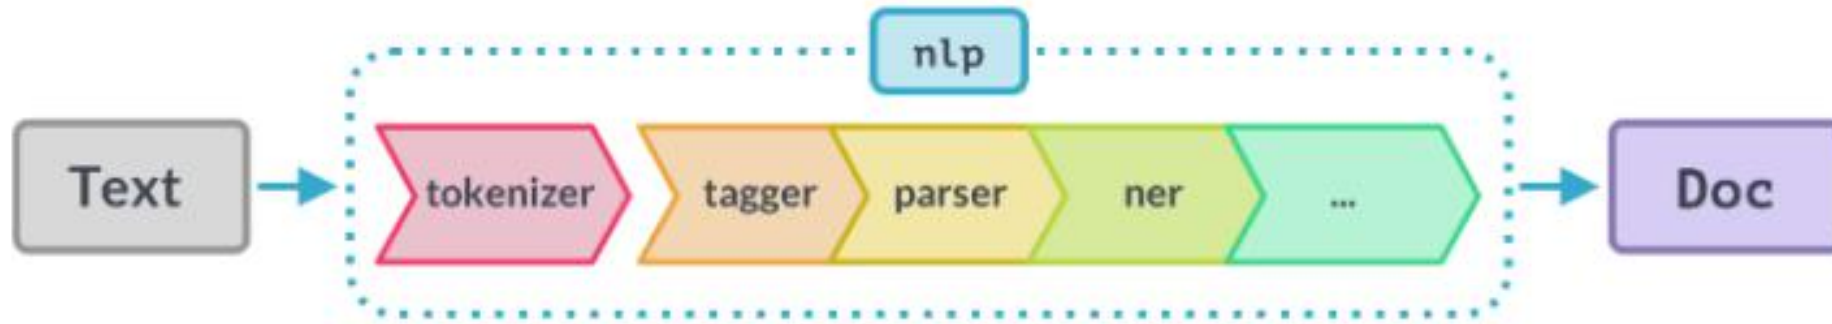
```
Doc.set_extension("page_number", default=None)
```

```
data = [  
    ("This is a text", {"id": 1, "page_number": 15}),  
    ("And another text", {"id": 2, "page_number": 16}),  
]
```

```
for doc, context in nlp.pipe(data, as_tuples=True):  
    doc._.id = context["id"]  
    doc._.page_number = context["page_number"]
```



## Using only the tokenizer (1)



- don't run the whole pipeline!

## Using only the tokenizer (2)

Use `nlp.make_doc` to turn a text into a Doc object

BAD:

```
doc = nlp("Hello world")
```

GOOD:

```
doc = nlp.make_doc("Hello world!")
```

## Disabling pipeline components

- Use `nlp.disable_pipes` to temporarily disable one or more pipes

# Disable tagger and parser

with `nlp.disable_pipes("tagger", "parser")`:

# Process the text and print the entities

`doc = nlp(text)`

`print(doc.ents)`

- Restores them after the with block
- Only runs the remaining components

## **Chapter 4: Training a neural network model**

# Training and updating models

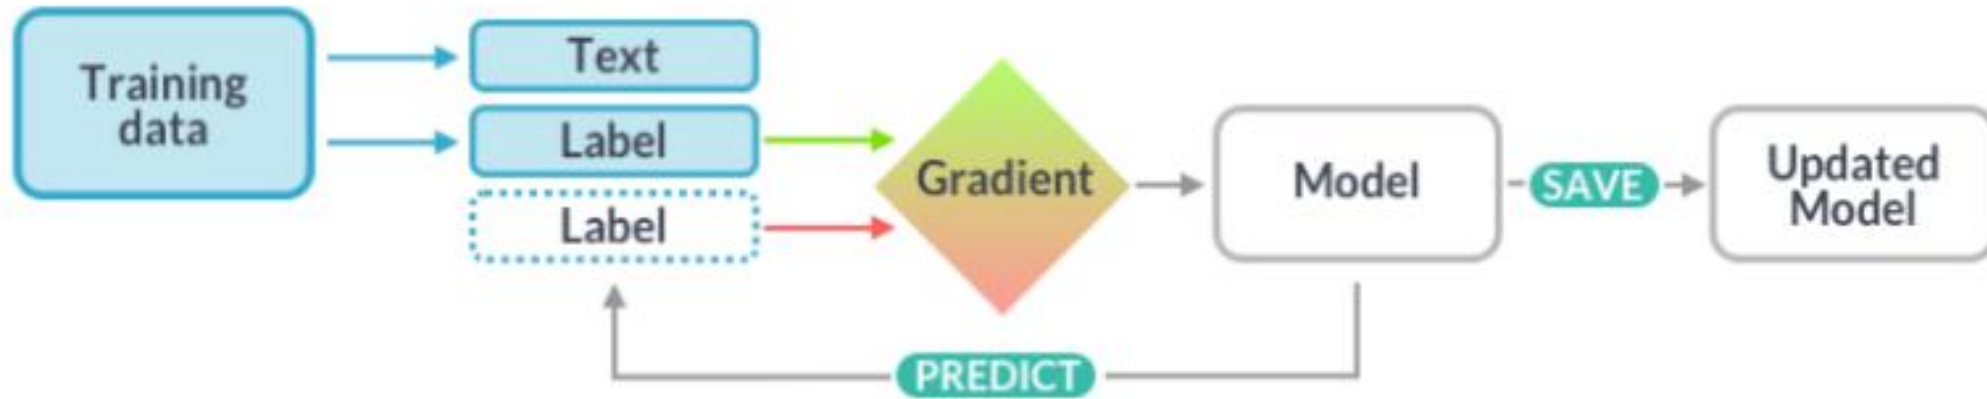
# Why updating the model?

- Better results on your specific domain
- Learn classification schemes specifically for your problem
- Essential for text classification
- Very useful for named entity recognition
- Less critical for part-of-speech tagging and dependency parsing

## How training works (1)

1. Initialize the model weights randomly with `nlp.begin_training`
2. Predict a few examples with the current weights by calling `nlp.update`
3. Compare prediction with true labels
4. Calculate how to change weights to improve predictions
5. Update weights slightly
6. Go back to 2.

## How training works (2)



- Training data: Examples and their annotations.
- Text: The input text the model should predict a label for.
- Label: The label the model should predict.
- Gradient: How to change the weights.



## Example: Training the entity recognizer

- The entity recognizer tags words and phrases in context
- Each token can only be part of one entity
- Examples need to come with context

("iPhone X is coming", {"entities": [(0, 8, "GADGET")]})

- Texts with no entities are also important

("I need a new phone! Any tips?", {"entities": []})

- Goal: teach the model to generalize

# The training data

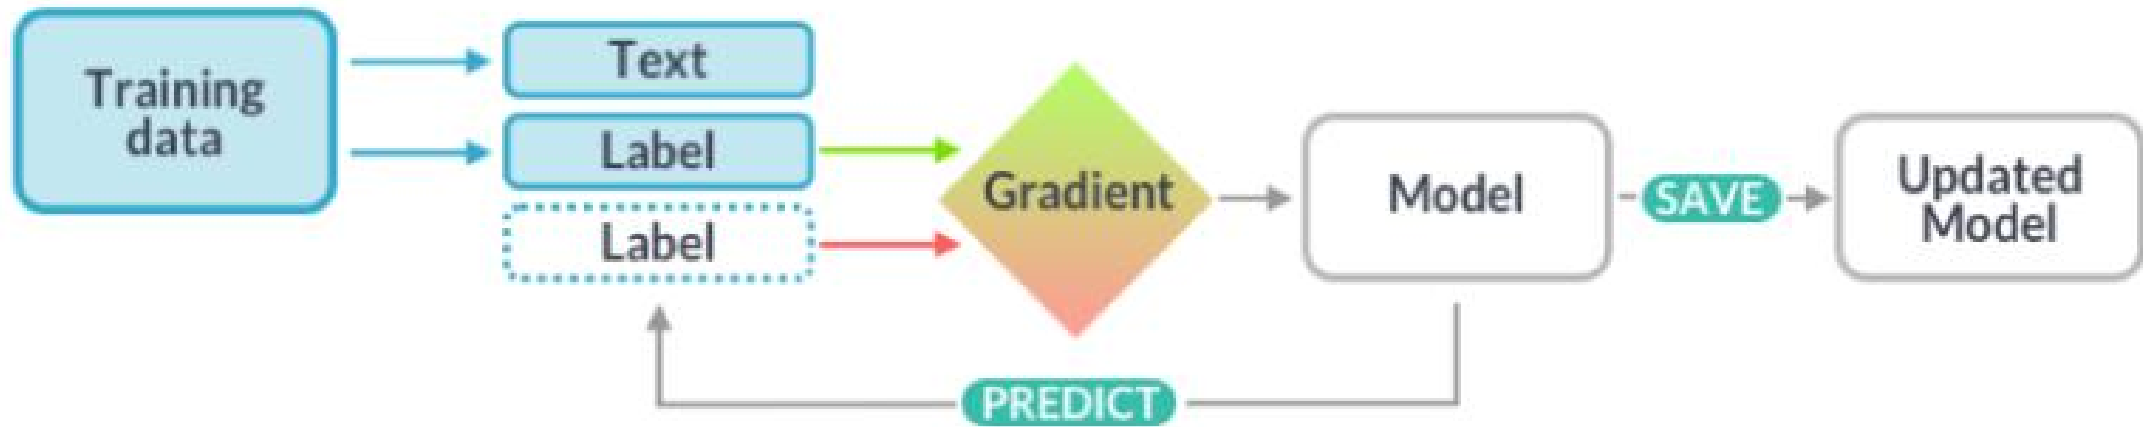
- Examples of what we want the model to predict in context
- Update an existing model: a few hundred to a few thousand examples
- Train a new category: a few thousand to a million examples
  - spaCy's English models: 2 million words
- Usually created manually by human annotators
- Can be semi-automated – for example, using spaCy's Matcher!

# The training loop

# The steps of a training loop

1. Loop for a number of times.
2. Shuffle the training data.
3. Divide the data into batches.
4. Update the model for each batch.
5. Save the updated model.

## Recap: How training works



- Training data: Examples and their annotations.
- Text: The input text the model should predict a label for.
- Label: The label the model should predict.
- Gradient: How to change the weights.

# Example loop

```
TRAINING_DATA = [  
    ("How to preorder the iPhone X", {"entities": [(20, 28, "GADGET")]}))  
    # And many more examples...  
]  
# Loop for 10 iterations  
for i in range(10):  
    # Shuffle the training data  
    random.shuffle(TRAINING_DATA)  
    # Create batches and iterate over them  
    for batch in spacy.util.minibatch(TRAINING_DATA):  
        # Split the batch in texts and annotations  
        texts = [text for text, annotation in batch]  
        annotations = [annotation for text, annotation in batch]  
        # Update the model  
        nlp.update(texts, annotations)  
  
# Save the model  
nlp.to_disk(path_to_model)
```

## Updating an existing model

- Improve the predictions on new data
- Especially useful to improve existing categories, like "PERSON"
- Also possible to add new categories
- Be careful and make sure the model doesn't "forget" the old ones

# Setting up a new pipeline from scratch

```
nlp = spacy.blank("en")# Start with blank English model
# Create blank entity recognizer and add it to the pipeline
ner = nlp.create_pipe("ner")
nlp.add_pipe(ner)
ner.add_label("GADGET")# Add a new label

# Start the training
nlp.begin_training()
# Train for 10 iterations
for itn in range(10):
    random.shuffle(examples)
    # Divide examples into batches
    for batch in spacy.util.minibatch(examples, size=2):
        texts = [text for text, annotation in batch]
        annotations = [annotation for text, annotation in batch]
        # Update the model
        nlp.update(texts, annotations)
```



# **Best practices for training spaCy models**

## Problem 1: Models can "forget" things

- Existing model can overfit on new data
  - e.g.: if you only update it with "WEBSITE", it can "unlearn" what a "PERSON" is
- Also known as "catastrophic forgetting" problem

## Solution 1: Mix in previously correct predictions

- For example, if you're training "WEBSITE", also include examples of "PERSON"
- Run existing spaCy model over data and extract all other relevant entities

BAD:

```
TRAINING_DATA = [  
    ("Reddit is a website", {"entities": [(0, 6, "WEBSITE")]}),  
]
```

GOOD:

```
TRAINING_DATA = [  
    ("Reddit is a website", {"entities": [(0, 6, "WEBSITE")]}),  
    ("Obama is a person", {"entities": [(0, 5, "PERSON")]}),  
]
```

## Problem 2: Models can't learn everything

- spaCy's models make predictions based on local context
- Model can struggle to learn if decision is difficult to make based on context
- Label scheme needs to be consistent and not too specific
  - For example: "CLOTHING" is better than "ADULT\_CLOTHING" and "CHILDRENS\_CLOTHING"

## **Solution 2: Plan your label scheme carefully**

- Pick categories that are reflected in local context
- More generic is better than too specific
- Use rules to go from generic labels to specific categories

**BAD:**

```
LABELS = ["ADULT_SHOES", "CHILDRENS_SHOES",  
"BANDS_I_LIKE"]
```

**GOOD:**

```
LABELS = ["CLOTHING", "BAND"]
```

## Wrapping up

## Your new spaCy skills

- Extract linguistic features: part-of-speech tags, dependencies, named entities
- Work with pre-trained statistical models
- Find words and phrases using Matcher and PhraseMatcher match rules
- Best practices for working with data structures Doc, Token Span, Vocab, Lexeme
- Find semantic similarities using word vectors
- Write custom pipeline components with extension attributes
- Scale up your spaCy pipelines and make them fast
- Create training data for spaCy's statistical models
- Train and update spaCy's neural network models with new data

## More things to do with spaCy (1)

- Training and updating other pipeline components
  - Part-of-speech tagger
  - Dependency parser
  - Text classifier



## More things to do with spaCy (2)

- Customizing the tokenizer
  - Adding rules and exceptions to split text differently
- Adding or improving support for other languages
  - 55+ languages currently
  - Lots of room for improvement and more languages
  - Allows training models for other languages